

## Exercise 13 – Multitasking and asyncio

### Objective

To run external programs, in this case other Python scripts, using a variety of methods, first using the **subprocess** module and then using **multiprocessing**.

### Questions

1. In the **labs** or (on Linux) your home directory, you will find a simple Python program, **client.py**, which lists files to STDOUT. The name of the file is specified at the command line, and if it cannot be read then an error is returned, using `exit`.
  - a) Now call the Python program **client.py** from another, passing a filename. If you can't think of a file to list, use the current program, or use the 'words' file.

Output an error message if, for some reason, the **client.py** fails.

Test this by:

- passing a non-existent file name.
  - calling a non-existent program.
- b) Modify the calling program to use a pipe and capture its output in a list. Print out the number of lines returned by the **client.py** program. Test as before.
  2. The purpose of this exercise is to experiment with different scenarios using the **multiprocessing** module. This is best demonstrated using a multi-core machine, so you might first like to check if that is the case. If not, then the exercise is still valid, but not quite so interesting.

**Note:** IDLE, and some other IDEs, does not display output from the child processes run by the multiprocessing module. So, run your code from the command-line.

Word prefixes are also called *stems*. We have written a program, **stems.py**, that reads the words file and generates the most popular stems of 2 to **n** characters long. It uses the **mytimer** module we created in a previous exercise, which you should make available.

Run the supplied **stems.py** program and note the time taken. You will note that no word exceeds 28 characters, so **n** could be 28. However, we can increase the value of **n** to obtain a longer runtime and demonstrate multiprocessing.

This time could be better used by splitting the task between cores. Using the **multiprocessing** module will require the stem search to be moved to a function. Make sure that all the rest of the code is only executed in main (if `__name__ == '__main__': test`).

Scenarios:

a) **n** worker processes.

This is where we split the task such that each stem length search runs in its own child process.

b) 2 worker processes **n/2** stem sizes each.

This assumes 2 CPU cores. It will require two processes to be launched explicitly, and each to be given a range of stem lengths to handle.

c) 2 worker processes using a queue.

This assumes 2 CPU cores. As in b), but instead of passing a range, pass the stem lengths through a queue. Make sure you have a protocol for the worker processes to detect that the queue has finished.

## Solutions

### Question 1

```
import subprocess
import os
import sys

#(a)
proc = subprocess.run([sys.executable, 'client.py', 'words'])
print('Child exited with', proc.returncode)

#(b)
proc = subprocess.run([sys.executable, 'client.py', 'words'],
                      stdout=subprocess.PIPE, stderr=subprocess.PIPE)

if proc.stderr != None:
    print('error:', proc.stderr.decode())

print('output:', proc.stdout.decode())
```

### Question 2

The timings will obviously vary depending on the machine:

```
a)
import mytimer
from multiprocessing import Process

def stem_search(stems, stem_size):
    best_stem = ""
    best_count = 0
    for (stem, count) in stems.items():
        if stem_size == len(stem) and count > best_count:
```

```
        best_stem = stem
        best_count = count
    if best_stem:
        print ('Most popular stem of size', stem_size, 'is:',
              best_stem, '(occurs', best_count, 'times)')
    return

if __name__ == '__main__':
    mytimer.start_timer()
    stems = {}
    for row in open('words', 'r'):
        for count in range(1, len(row)):
            stem = row[0:count]
            if stem in stems:
                stems[stem] += 1
            else:
                stems[stem] = 1
    mytimer.end_timer('Load')

    # Process the stems.
    mytimer.start_timer()
    n = 30
    for stem_size in range(2, n+1):
        proc = Process(target=stem_search,
                      args=(stems, stem_size))
        proc.start()
        processes.append(proc)
    for proc in processes:
        proc.join()
        mytimer.end_timer('Process')
```

b)

```
import mytimer
from multiprocessing import Process

def stem_search(stems, start, end):
    for stem_size in range(start, end):
        best_stem = ""
        best_count = 0
        for (stem, count) in stems.items():
            if stem_size == len(stem) and
                count > best_count:
                best_stem = stem
                best_count = count

        if best_stem:
            print ('Most popular stem of size',
                    stem_size, 'is:', best_stem,
                    '(occurs', best_count, 'times)')

    return

if __name__ == '__main__':
    mytimer.start_timer()
    stems = {}
    for row in open('words', 'r'):
        for count in range(1, len(row)):
            stem = row[0:count]
            if stem in stems:
                stems[stem] += 1
            else:
                stems[stem] = 1
    mytimer.end_timer('Load')
    # Process the stems.
```

```
mytimer.start_timer()
n = 30
proc1 = Process(target=stem_search,
                 args=(stems, 2, int(n/2) + 1))
proc1.start()
proc2 = Process(target=stem_search,
                 args=(stems, int(n/2) + 1, n + 1))
proc2.start()
proc1.join()
proc2.join()
mytimer.end_timer('Process')
```

c)

```
import mytimer
from multiprocessing import Process, Queue

def stem_search(stems, queue):
    stem_size = 1
    while stem_size > 0:
        stem_size = queue.get()
        best_stem = ""
        best_count = 0

        for (stem, count) in stems.items():
            if stem_size == len(stem) and count > best_count:
                best_stem = stem
                best_count = count

        if best_stem:
            print ('Most popular stem of size', stem_size,
                  'is:', best_stem, '(occurs', best_count,
```

```
        'times)')
        return

if __name__ == '__main__':
    mytimer.start_timer()
    stems = {}
    for row in open('words', 'r'):
        for count in range(1, len(row)):
            stem = row[0:count]
            if stem in stems:
                stems[stem] += 1
            else:
                stems[stem] = 1
    mytimer.end_timer('Load')
    mytimer.start_timer()
    n = 30
    queue = Queue()
    proc1 = Process(target=stem_search, args=(stems, queue))
    proc2 = Process(target=stem_search, args=(stems, queue))
    proc1.start()
    proc2.start()
    for stem_size in range(2, n):
        queue.put(stem_size)
    queue.put(0)
    queue.put(0)
    proc1.join()
    proc2.join()
    mytimer.end_timer('Process')
```